elucidata_demonstrator_3_3

November 18, 2022

```
[1]: import support as sp
     import visualization as vis

     %load_ext autoreload
     %autoreload 2
```

1 Starter Kit 3.3: Feature Engineering

1.1 Description

Most data mining and machine learning algorithms do not work well if you just feed them your raw data: such data often contains noise and the most relevant distinguishing information is often implicit. For example, raw sensor data just contains a timestamp and a corresponding value, while interesting aspects of it might be the trends contained within or the number of times a threshold is exceeded.

Feature engineering is the process of extracting from the raw data the most relevant distinguishing characteristics that will be presented to the algorithm during modeling. It is a way of deriving new information from the existing data, so that its characteristics are more explicitly represented. The resulting features are eventually feed to the AI/ML algorithm of the model. In practice, the feature engineering step is achieved by the selection of the most appropriate parameters, or the composition of new features by manipulation, transformation and combination of the raw data.

Feature engineering is one of the most important and creative steps in the data science workflow. However, there is no clearly-defined formal process for engineering features and, consequently, this requires a lot of creativity, a good understanding of the domain and of the available data, some trial-and-error, etc. It is also desirable to have upfront an idea of the modeling and analysis task for which you want to use the resulting features, as this might help you to identify relevant features.

1.2 Business goal

The overall goal of this Starter Kit is to present some **advanced feature engineering steps** related to **feature construction** and **extraction**. By keeping in mind **what is your business question** and **what is the corresponding data science task**, you will be able to derive **valuable features** that can be used in the next stage of your analysis.

1.3 Application context

Feature engineering is one of the steps in the data science workflow with the most decisive impact on the accuracy of the model you want to develop. It is the final step before actually training the model and it defines how the input data will be fed to the model.

1.4 Starter Kit outline

In this Starter Kit we use a dataset from the Regional Rail System of Pennsylvania. Before starting the feature extraction *per se*, we will first apply some basic preprocessing to the dataset and have a high-level overview of the data. We will then derive several interesting features from the dataset that can be used to characterise train rides and delays.

1.5 Basic data understanding

In order to illustrate how to engineer features, we will use in this Starter Kit a dataset provided by the Southeastern Pennsylvania Transportation Authority (SEPTA), which can be downloaded here.

The SEPTA Regional Rail System consists of 13 branches and more than 150 stations in Philadelphia, Pennsylvania, and its suburbs and satellite cities, as depicted in the map below.

SEPTA uses On-Time Performance (OTP) to measure service reliability. OTP identifies the number of trains for all rail lines that arrive at their scheduled destination at the scheduled time. However, by industry standard, a train may arrive up to 5 minutes and 59 seconds after its scheduled time and still be considered on-time.

SEPTA has established an annual goal of 91% for Regional Rail On-Time Performance. Based on this goal, we will demonstrate how feature engineering can be used to validate SEPTA's requirement. More specifically, we will start designing features that may help us **predict** delays of a train on a given **stop** at a given **date**.

The SEPTA dataset consists of 2 CSV files that we will consider in this Starter Kit. The first file, `otp.csv`, provides OTP data, that is, information about trains and the times at which they arrive at the different stations on their lines. The table below shows an excerpt of this OTP dataset.

```
[2]: df_otp = sp.load_otp_data()
     df_otp.head(10)
```

```
2022-11-18 14:24:53,784 [INFO ]  Using file: /Users/mdhn/Documents/starter-
kits/repository/elucidatalab.starterkits/data/SK_3_3/otp.csv
```

```
[2]:   train_id direction                 origin          next_station        date  \
     0       778         N               Trenton               Stenton  2016-03-23
     1       598         N             Thorndale              Narberth  2016-03-23
     2       279         S                   Elm           Ridley Park  2016-03-23
     3       476         N  Airport Terminal E-F      Suburban Station  2016-03-23
     4       474         N  Airport Terminal E-F   Jenkintown-Wyncote  2016-03-23
     5       279         S                   Elm            Crum Lynne  2016-03-23
     6       778         N               Trenton              Sedgwick  2016-03-23
     7       598         N             Thorndale                Merion  2016-03-23
```

```
8     395         S                   Trent              Temple U 2016-03-23
9    6464         N          Powelton Yard  30th Street Station 2016-03-23

     status            timeStamp
0    1 min 2016-03-23 00:01:47
1    1 min 2016-03-23 00:01:58
2    2 min 2016-03-23 00:02:02
3  On Time 2016-03-23 00:03:19
4  On Time 2016-03-23 00:03:35
5    2 min 2016-03-23 00:03:42
6    2 min 2016-03-23 00:03:48
7    1 min 2016-03-23 00:03:58
8  On Time 2016-03-23 00:04:17
9  On Time 2016-03-23 00:04:22
```

The different attributes in this file OTP dataset are as follows: * `train_id`: the identifier of the train * `direction`: the train direction; its values are 'N' for Northbound and 'S' for Southbound * `origin`: the station before *next_station* * `next_station`: the next station stop at *timeStamp* * `date`: the date of the journey * `status`: the current train delay; its values are 'On Time' or the delay, i.e. the amount of time above the 5min59s limit until which a train is considered to be on time (e.g. '1 min', '5 min', '10 min'); a value of 999 indicates a suspended train * `timeStamp`: the timestamp at which the train will arrive at `next_station`

The second file, `trainView.csv`, provides train tracking information:

```
[3]: df_train_view = sp.load_train_view_data()
     df_train_view.head(3)
```

2022-11-18 14:25:08,389 [INFO] Using file: /Users/mdhn/Documents/starter-
kits/repository/elucidatalab.starterkits/data/SK_3_3/trainView.csv

```
[3]:   train_id status next_station service            dest       lon       lat  \
     0    102TT      0       Radnor   LOCAL  Colmar-Link Belt -75.37250  40.04388
     1    102TT      0   St. Davids   LOCAL  Colmar-Link Belt -75.38670  40.04583
     2    102TT      0    Strafford   LOCAL  Colmar-Link Belt -75.42277  40.04722

       source track_change track       date            timeStamp0  \
     0  Devon           -1    -1 2016-04-22  2016-04-22 13:21:07
     1  Devon           -1    -1 2016-04-22  2016-04-22 13:19:11
     2  Devon           -1    -1 2016-04-22  2016-04-22 13:15:04

               timeStamp1  seconds
     0  2016-04-22 13:22:43       96
     1  2016-04-22 13:21:01      110
     2  2016-04-22 13:17:01      117
```

Its most important attributes are: * `lon`: the current GPS longitude of the train * `lat`: the current GPS latitude of the train * `timeStamp0`: the earliest timestamp at the current GPS coordinates *

timeStamp1: the latest timeStamp at the current GPS coordinates * `seconds`: the time difference (in seconds) between `timeStamp1` and `timeStamp0` * `track_change`: The name of the track if there was a track change, else -1

We can already observe in the previews of these two datasets format differences at the level of the `train_id` and `status` columns. We will address these in the next section.

1.6 Data preprocessing

1.6.1 Removing invalid data instances

A manual inspection of the dataset reveals that the `train_id` column from the OTP data contains 1 negative value and many values containing letters and punctuation characters. Let's filter out the rows for which the `train_id` is not a number.

```
[4]: df_otp = sp.remove_invalid_data_instances(df=df_otp)
```

```
We found 35964 IDs with non-numeric characters, which is 1.91% from a total of
1882015 IDs.
In addition, there are 88333 entries with next_station name = None
These entries with an unexpected train ID or no station name information will be
excluded from the dataset
```

1.6.2 Understanding train_id

`train_id` can be used an identifier for a specific train run between a departure and a destination station at a given time of the day. To confirm that this is possible, we count how many times a `train_id` passes through a given station on a given day in the `train_run` dataset:

```
[5]: train_run = sp.get_train_run(df=df_otp)
     print("In %0.2f%% of the cases, a train_id passes only once through a station␣
       ↪on a given day." %
             (train_run.single_pass.sum() / float(len(train_run)) * 100))
```

```
In 99.81% of the cases, a train_id passes only once through a station on a given
day.
```

We will exclude from the dataset the remaining 0.2% of cases, since they correspond to exceptional situations where a `train_id` was registered more than once in a day. For consistency, we only keep in the trainView dataset those rows for which the (`train_id`, `next_station`, `date`) tuples also occur in the OTP dataset.

For the remainder of the notebook, we will refer to **train id** as an identifier for a train journey between a given *origin* and a given *destination* at a given *time* of the day. In other words, a train id can be repeated on multiple days. A **train run** on the other hand, is now defined as a train id on a specific date, as used in the trainview dataset.

```
[6]: df_otp, df_train_view = sp.exclude_remaining_cases(otp=df_otp,␣
       ↪train_view=df_train_view)
```

1.6.3 Turning statuses into workable delays

The current string-based format of the `status` column in the OTP dataset is not suitable for performing calculations on the delays. We will transform its values into integers (in minutes) and rename the column to `delay`.

```
[7]: df_otp = sp.calculate_delays_for_otp(df=df_otp)
```

```
      status  delay
0     1 min       1
1     1 min       1
2     2 min       2
3   On Time       0
4   On Time       0
```

We will also exclude missing (`None`) values present in the `status` column of the df_train_view dataset, and, as for the OTP dataset, convert the values of that column to integers and rename the column to `delay`.

```
[8]: df_train_view = sp.calculate_delays_for_train_view(df=df_train_view)
```

1.6.4 Removing suspended trains

Since a status of `999` represented a suspended train, we exclude those trains from the datasets. In addition, a delay of 1440 minutes corresponds to one day delay, meaning that the train was most likely canceled. We exclude the trains with those delays from the datasets as well.

```
[9]: df_otp, df_train_view = sp.remove_suspended_trains(df_otp=df_otp,␣
     ↪df_train_view=df_train_view)
```

1.7 Data overview

Let's check some basic characteristics of the (cleaned OTP) dataset.

```
[10]: sp.get_otp_overview(df=df_otp)
```

```
[10]:
      Number of stations                              154
      Number of train ids                             995
      Number of train runs                         173924
      Time range              23 Mar 2016 to 06 Nov 2016
      Percentage trains on time                      74.9%
```

As the table above indicates 74.9% of the trains being delayed, i.e. we demonstrate analytically that **the original requirement of having 91% of all trains on time** (definition: no more than 5min59s delay) **is far from being reached**. This is just a showcase of how feature engineering can be used in practice. In the remainder of the notebook we will present further examples to better understand the process.

1.8 Feature engineering

We are now ready to investigate several features to demonstrate the workflow and some hidden obstacles while calculating features. You might come up with other features as well, so don't hesitate to try them out.

Based on our interest to characterize train rides and delays, we come up with the following features: - Day of the week - Month of the year - Rush hours - Rank the different stops in a train run - Distance between stations - Cumulative distance along a train run - Total distance of a train run - Northbound vs Southbound travels - Stations with big delay - Delay over the last 7 days - Track changes

1.8.1 Day of the week

Based on our collective experience from using public transport, we can expect that delays are more likely to happen at specific moments in a day (e.g. rush hours) and at specific days in a week (e.g. working days). In order to be able to easily verify these assumptions later on, we first extract for each entry in the OTP dataset the following features: - the hour - the name of the day - the type of day, which we define as having 3 possible values: weekday, Saturday and Sunday - and a boolean indicating whether that day is a weekday or a weekend day

Let's start by looking at the impact of day of the week on delay.

```
[11]: vis.plot_train_delays(df=df_otp)
```


As we can see, the difference in delays on weekdays (Mon-Fri) is only minimal. We will thus

categorize the day of the week into weekdays, saturdays and sundays.

```
[12]: df_otp = sp.get_type_of_days(df=df_otp)
```

1.8.2 Month of the year

Seasonal effects, expressed here as the month of the year, can have an impact on train performance. As an example consider the summer holidays one can expect less (commuting) traffic, albeit also reduced staff, which might affect how on-time trains are. Additionally, one might expect more (bad) weather related problems with trains or tracks during winter. Let's start by looking at delay as a function of the month of the year.

```
[13]: vis.plot_train_delays(df=df_otp, kind='month')
```

Train delay distribution per month

Data only runs from March to November 2016. As we can see, there is a large variability observed between months. As a consequence, we will keep this feature, rather than grouping the months into the four seasons.

1.8.3 Rush hours

During rush hours there's an intense increase in the number of passengers and, often, of the number of trains. It is also a period where it is most critical for trains to be on-time. To identify rush hour periods, we will start by extracting some time features from the dataset: the hour and whether it is a weekend day or not.

```
[14]: df_otp = sp.extract_time_features(df=df_otp)
      df_otp[['train_id', 'origin', 'date', 'hour', 'isWeekend']].sample(5)
```

```
[14]:          train_id              origin        date hour  isWeekend
      1695170      5313          Doylestown  2016-10-29   11       True
      147377       7843   Chestnut Hill East  2016-04-08   17      False
      70165         470  Airport Terminal E-F  2016-03-30   22      False
      378319        519            Tabor Jct  2016-05-05   09      False
      1414925      5325            Lansdale  2016-09-25   17       True
```

Let's now use these features to count the number of distinct trains in the rail system per hour during weekdays, and hence have an idea of the train density per hour. We define a rush hour as one for which the number of trains per hour and across all hours and weekdays is above the 75th percentile. Let's visualize the train density per hour on weekdays with that threshold.

```
[15]: vis.plot_train_density(df=df_otp)
```


# of distinct trains per hour across all weekdays (Mon-Fri)

8

We will now define rush hour as all times on weekdays that go beyond the 75th percentile, which is 135 trains per hour as indicated with the red line in the plot above. As an example, we will also provide an excerpt of the OTP data.

```
[16]: df_otp = sp.define_rush_hour(df=df_otp, percentile=75)
      df_otp[['train_id', 'origin', 'date', 'hour', 'dayOfTheWeek', 'isRushHour']].
        ↪sample(5).query("isRushHour==True").head()
```

The threshold that we will use is of 135 trains per hour over all weekdays and
hours.
Identified rush hours are: 07, 08, 09, 16, 17, 18, 19.

```
[16]:        train_id origin       date hour dayOfTheWeek  isRushHour
      49655       388  Elwyn 2016-03-28   19          Mon        True
```

1.8.4 Rank the different stops in a train run

The fact that a `train_id` passes only once through a station on a given day allows us to assume that a `train_id` is defined for a specific trip (origin-destination) at a given time in the day. So the combination of a `train_id` and `date` identifies a unique train run (as we defined before). We construct the feature *rank* to identify the sequence number of each stop along the train run, i.e., 1 is origin station, 2 is the second stop, 3 is the third stop, etc.

```
[17]: df_otp = sp.calculate_rank_stop(df=df_otp)
```

We can now look at the delay as a function of the rank number.

```
[18]: vis.plot_delays_per_stop_number(df=df_otp)
```


9
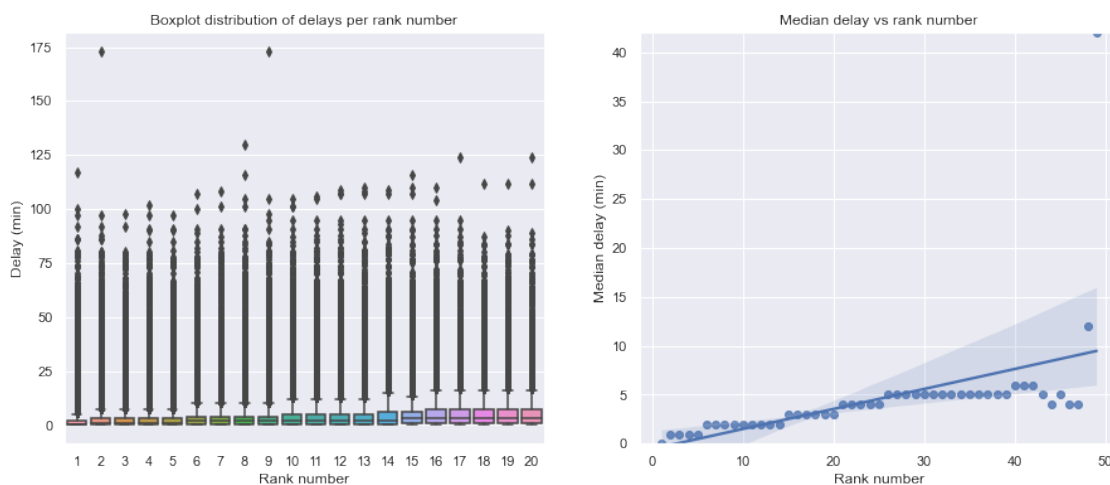
We can see that the delay is a function of the train stop number, confirming the idea that trains accumulate delays along the run, without being able of (fully) compensating for them.

1.8.5 Distance between stations

We will now use the coordinates information to derive the distance in km between stations, the total distance per train run and the cumulative distance along a train run. Since there are multiple coordinate points for each stop (train position might be registered at different time points, e.g. when entering the station, when docking, when departing, etc.), we will take the average latitude and longitude for each station.

```
[19]: uni_coords, uni_runs, df_distance, df_coords = sp.
      ↪calculate_distance_between_stations(df_otp=df_otp,␣
      ↪df_train_view=df_train_view)
      df_coords[['train_id', 'date', 'rank', 'next_station', 'upcoming_station',␣
      ↪'distance']].sample(5)
```

```
[19]:          train_id        date  rank next_station upcoming_station  distance
      353974       9346  2016-06-09   2.0        Media    Elwyn Station  0.257473
      111792        313  2016-05-12  16.0    Wayne Jct     Fern Rock TC  2.506551
      762170       3522  2016-08-10   7.0     Fernwood        Lansdowne  1.379204
      1410438      3534  2016-08-02  30.0     Lansdale        Pennbrook  3.891134
      761099       3240  2016-07-01   6.0     Fernwood        Lansdowne  1.379204
```

We will now visualize an example of a train run. We will plot the locations of all the train stations (in gray) with the stations of the example train run colored by the order they occur.

```
[20]: vis.plot_train_run(df_train_view=df_train_view, uni_coords=uni_coords,
                         uni_runs=uni_runs, df_distance=df_distance,␣
      ↪df_coords=df_coords)
```

If we plot the distribution of distances between consecutive train stations we can see that most inter-station distances are shorter than 10km.

```
[21]: vis.plot_distance_distribution(df=df_coords)
```

1.8.6 Cumulative distance along a train run

This feature provides the cumulative distance that a train has traveled along a given train run. It is likely correlated with the rank of a given station on a given train run. Nevertheless, it might convey different information and be useful, depending on the specific questions we will try to address, when applying this feature engineering analysis to a machine learning task.

In the table below we see the information regarding distance and cumulative distance for an example train run.

```
[22]: df_coords, random_train_run = sp.
      ↪calculate_cumulative_distance_along_train_run(df=df_coords,␣
      ↪uni_runs=uni_runs)
      random_train_run
```

```
[22]:     train_id  rank        next_station  delay     upcoming_station  distance  \
      0        446   2.0   Airport Terminal B      0  Airport Terminal C-D  0.156816
      1        446   3.0      Eastwick Station      0    Airport Terminal B  4.695298
      2        446   4.0       University City      0      Eastwick Station  5.474109
      3        446   5.0   30th Street Station      0       University City  1.912928
      4        446   6.0       Suburban Station      0   30th Street Station  1.567499
      5        446   7.0      Jefferson Station      0      Suburban Station  1.344006
      6        446   8.0              Temple U      0     Jefferson Station  3.042680
      7        446   9.0             Wayne Jct      0              Temple U  2.582294
      8        446  10.0          Fern Rock TC      0             Wayne Jct  2.506551
      9        446  11.0          Melrose Park      0          Fern Rock TC  2.424309
      10       446  12.0           Elkins Park      0          Melrose Park  3.024633
      11       446  13.0   Jenkintown-Wyncote      1           Elkins Park  0.524916
      12       446  14.0              Glenside      1   Jenkintown-Wyncote  1.991849
      13       446  15.0               Ardsley      1              Glenside  1.410460
      14       446  16.0                Roslyn      0               Ardsley  2.039445
      15       446  17.0             Crestmont      1                Roslyn  0.867242
      16       446  18.0           Willow Grove      2             Crestmont  2.326806
      17       446  19.0               Hatboro      0          Willow Grove  9.882850
      18       446  20.0            Warminster      0               Hatboro  5.360457

          cum_distance
      0        0.156816
      1        4.852114
      2       10.326223
      3       12.239151
      4       13.806650
      5       15.150656
      6       18.193336
      7       20.775630
```

```
8       23.282180
9       25.706489
10      28.731122
11      29.256038
12      31.247886
13      32.658347
14      34.697791
15      35.565034
16      37.891839
17      47.774689
18      53.135146
```

In the scatter plot below we plot the delay on both the distance between two stops and the cumulative distances within the train runs. To be able to detect a correlation between those features, we plot the regression line in red on top of the scatter plot.

```
[67]:  vis.delay_jointplot(df=df_coords, uni_runs=uni_runs)
```

Delay as a function of distance between stations (limited to distances < 10km)

Delay as a function of cumulative distance of train run

Interestingly, only cumulative distance appears to show a correlation with delay (indicated by the positive regression line). The latter is suggesting that a delay might not be caused by the distance between stations, but rather by delays incurred at stations themselves. However, the longer the train runs, the longer will be the delay incurred (as we have already shown with stop rank).

1.8.7 Total distance of a train run

The total distance traveled by the train is probably closely linked to the total delay or the delay at the last stop of a train run. Here we will look at train run total distance and delay.

In the first scatter plot below we plot the total delay as a function to the total train run distance. Additionally, we plot the delay at the last stop as a function to the total train run distance. Here again, the red line indicates the regression line.

```
[68]: vis.total_train_run_distance_joint_plot(df=df_coords)
```

Total train run delay as a function total run distance

Delay at last stop as a function total run distance

We can see that train run distance clearly influences the total delay throughout the run, as well as the delay at the last stop.

```
[25]: df_otp = sp.add_cum_delay(df=df_otp)
      df_otp = sp.add_distance(df_otp=df_otp, df_coords=df_coords)
```

1.8.8   Northbound vs Southbound travels

This feature identifies stations which are typically problematic, in the sense that there are always delays. We can do this by checking their average delay.

We will first compare the delays of trains operating in one or the other direction.

```
[26]: vis.plot_northbound_southbound_dely_per_station(df=df_otp)
```

Delay per station for Northbound and Southbound journeys
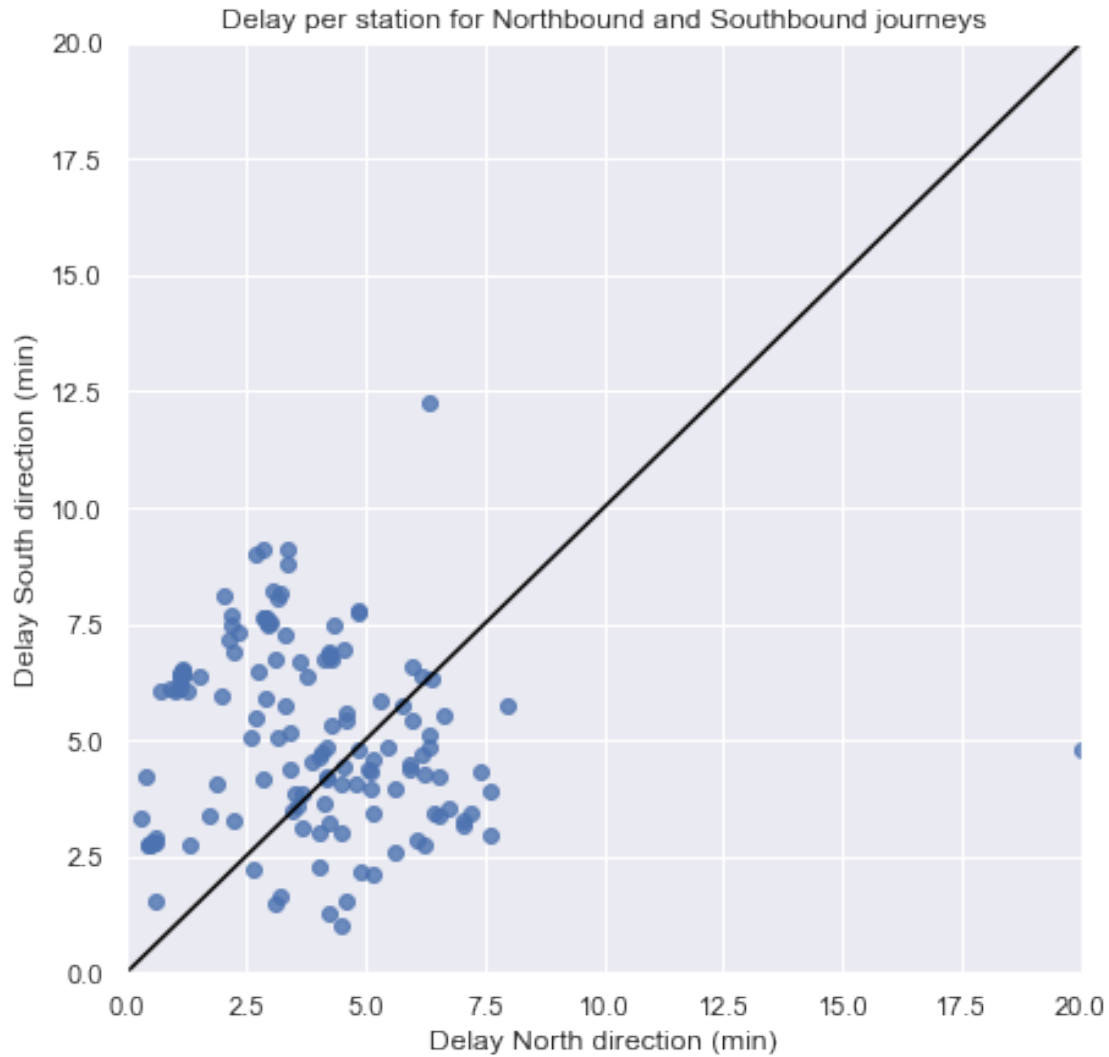
The asymmetry in the scatterplot (even if not that striking) suggests that delay is dependent on direction (as could be expected, since the distance from origin and stop number will be different according to the direction), so we will label long-duration stations per direction.

1.8.9 Stations with big delay

Some stations might be more prone to delays. This might be caused e.g. by many tracks converging to a smaller number of platforms of the station or due to long-term works on a station. We will, thus, rank the average delay for each station and label the top and bottom ones.

```
[27]: vis.plot_stations_with_long_delays(df=df_otp)
```

Direction: N

Direction: S

We can label those stations with an average delay above a certain threshold. In this case, we will use the 90th percentile as a threshold. Below you can see the histogram of the average delay per station.

```
[28]: vis.plot_average_delay(df=df_otp)

      sp.print_delays_overview(df=df_otp)
      df_otp = sp.label_long_delay_stations(df=df_otp)
```

Stations with an average delay above 7.4 minutes will be labeled:
        Direction N: Churchmans Crossing, Conshohocken, Elm St, Main St,

```
Norristown TC, Spring Mill
        Direction S: 49th St, Angora, Berwyn, Clifton-Aldan, Daylesford, Devon,
Downingtown, Exton, Fernwood, Malvern, Morton, Moylan-Rose Valley, Primos,
Radnor, Rosemont, Secane, St. Davids, Strafford, Swarthmore, Thorndale,
Villanova, Wayne-A, West Trenton, Whitford
```


Average station delay

The red line on the histogram above indicates the threshold (90th percentile) that was used to label stations with considerably long delays. Most of those stations are in direction S (southbound).

1.8.10 Delay over the last 7 days

Sometimes, uncontrollable factors can cause delays in trains for a couple of days (e.g. works on the train tracks, long duration strikes, etc.). We can calculate what was the delay over the last 7 days of each date to add that as a feature to predict delays.

```
[29]: df_otp, last_week_delay = sp.calculate_longest_delays_over_past_week(df=df_otp)
      print('10 cases with the longest delay over the past 7 days')
      last_week_delay.sort_values('last_week_delay', ascending=False).head(10)
```

```
10 cases with the longest delay over the past 7 days
```

```
[29]:       train_id          next_station  last_week_delay        date
      6278      6803  30th Street Station            116.0  2016-04-01
      7028       852  30th Street Station             88.0  2016-11-02
      1716       393              Yardley             86.0  2016-11-02
      315       1556     Jefferson Station             85.0  2016-09-17
      128       1524               Berwyn             85.0  2016-07-19
      129       1524           Daylesford             85.0  2016-07-19
```

```
130        1524              Devon       85.0 2016-07-19
134        1524            Malvern       85.0 2016-07-19
138        1524              Paoli       85.0 2016-07-19
139        1524             Radnor       85.0 2016-07-19
```

1.8.11   Track changes

Track changes can be the cause of delays. To address that, we will compare delays in cases when there was a track change with cases when there was no track changes.

```
[30]: df_track_changes = sp.get_track_changes(df=df_train_view)
      vis.plot_track_changes(df=df_track_changes)
```
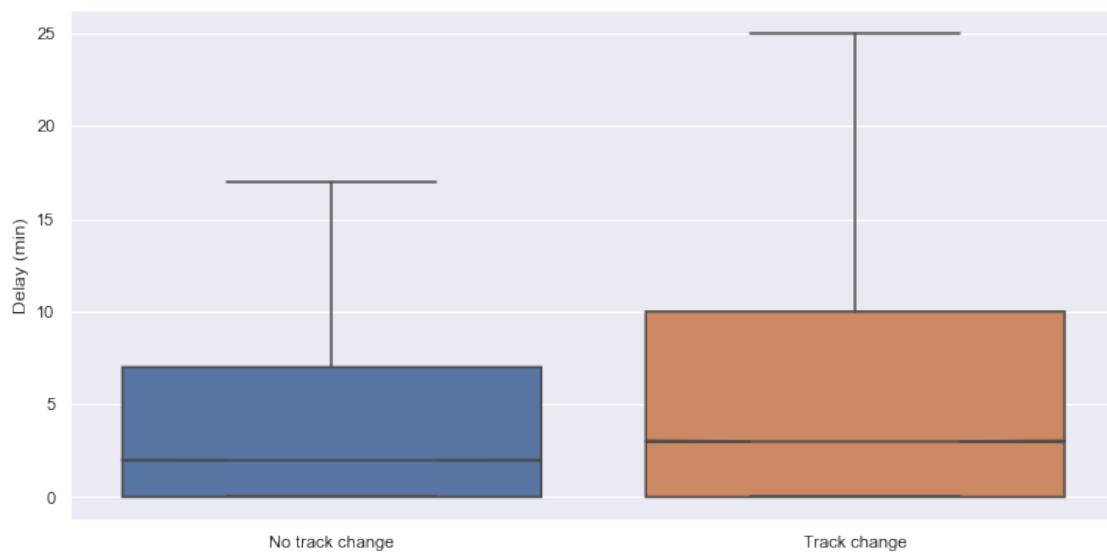

As we can see, delays are higher when there were track changes. A feature we can add to our model is the track change frequency per station per day, defined as:

$$Track change frequency_d = \frac{\sum_{tc\_tr}}{\sum_{tr}}$$

where d is a given date, tc_tr is a train arrival with a track change and tr is a train.

```
[31]: track_change_frequency_date = sp.
        ↪calculate_track_changes_frequency(df=df_track_changes)
      track_change_frequency = sp.
        ↪calculate_aggregated_track_changes(df=track_change_frequency_date)
```

We can compare how delay changes as a function of track change frequency:

```
[69]:
```

20

```
vis.
  ↪plot_track_change_frequency_with_delays(df_track_change_frequency=track_change_frequency_da
  ↪df_otp=df_otp)
```

Pearson's coefficient = 0.16 with p-value = 5.70e-178.

## Track change frequency as a function of the mean delay
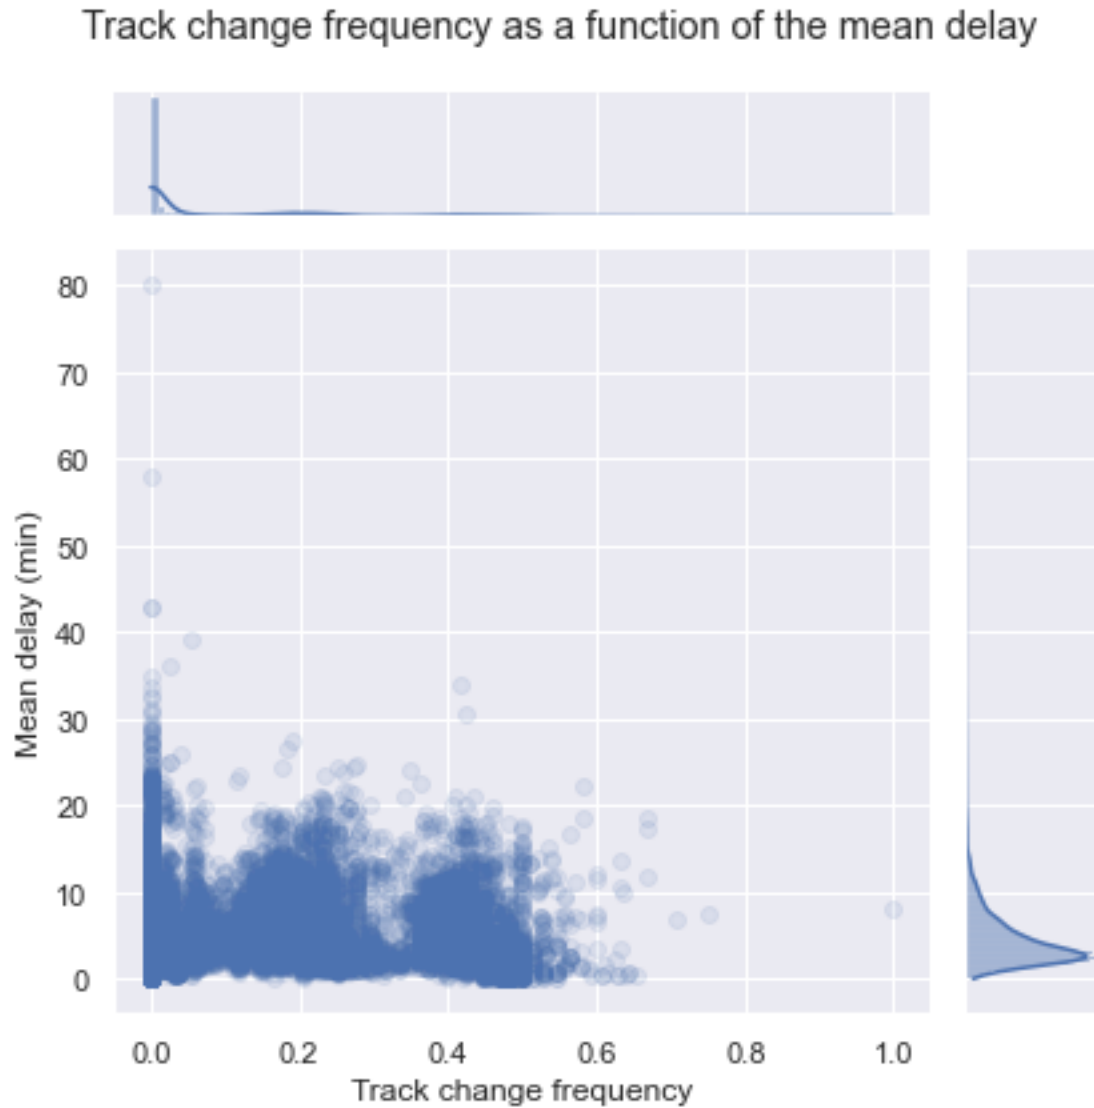
The red line indicates an upward trend, which is confirmed by the significant (i.e. the p-value >
0.05) positive Pearson correlation coefficient. The latter is suggesting that stations with higher
delays incur more track changes.

In the plot below we see the track change frequency of the 10 stations with the highest track change
frequency.

```
[33]:  # get unique combinations of station/coordinates and of train_id/dates
       uni_coords = df_train_view[['next_station', 'lon', 'lat']].drop_duplicates()
       uni_runs = df_otp[['train_id', 'date']].drop_duplicates()

       vis.
        ↪plot_highest_track_change_frequency(track_change_frequency=track_change_frequency)
```


In the plot below you can see the same figure, but now for the 10 stations with the lowest change frequency. Note that the change frequency for station Cynwyd is zero, which is enforced by the physical limitation of only a single track.

```
[34]:  vis.plot_lowest_track_change_frequency(track_change_frequency)
```

```
[35]: # add track change frequency as a feature
      df_otp = df_otp.merge(track_change_frequency, on='next_station')
```

1.9 Conclusion

In this Starter Kit we demonstrated the workflow on how to construct advanced features that are fundamental for answering company business questions such as SEPTA' performance measures. We demonstrated the need of data preparation steps before calculating features as it might simplify the work later. In addition, a wide range of complementary features are constructed, while the impact on train delay is briefly examined using statistical plots.